

Lab 3 - Jackson Nahom

September 15, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn import metrics
```

```
[2]: # load txt file
names = pd.read_csv('data/kddcup.names.txt', header=None, delimiter=':
↳', skiprows=1)

# make column 0 into a list
name_list = names[0].tolist()

# add the last column with type
name_list.append('type')
```

```
[3]: # use the column names
netattacks = pd.read_csv('data/kddcup.data_10_percent.gz', names=name_list,
↳ header=None, index_col=None)

# netattacks.head()
# netattacks.describe(include='all')
```

```
[4]: netattacks['label'] = np.where(netattacks['type'] == 'normal.', 'good', 'bad')
# netattacks['label'].value_counts()
```

```
[5]: train, test = train_test_split(netattacks, test_size=0.25)
print("Rows in train:", len(train))
print("Rows in test:", len(test))
```

Rows in train: 370515

Rows in test: 123506

```
[6]: # define new tree
dt = tree.DecisionTreeClassifier()
# train the model using a list of column names
pred_vars = ['duration', 'src_bytes', 'dst_bytes']
```

```
# The value we are trying to predict is 'label'
dt.fit(train.loc[:, pred_vars], train['label'])
# tree.plot_tree(dt)
```

```
[6]: DecisionTreeClassifier()
```

```
[7]: predicted = dt.predict(test.loc[:, pred_vars])
      print(predicted[:5]) # show first five predictions
```

```
['bad' 'bad' 'bad' 'bad' 'bad']
```

```
[8]: from collections import Counter
      # count test data
      test_labels_stats = Counter(test['label'])
      print("Labels in the test data:", test_labels_stats)

      # count predicted
      predicted_labels_stats = Counter(predicted)
      print("Labels in the predictions:", predicted_labels_stats)
```


```
Labels in the test data: Counter({'bad': 99028, 'good': 24478})
```

```
Labels in the predictions: Counter({'bad': 100403, 'good': 23103})
```

```
[9]: metrics.confusion_matrix(y_true=test['label'], y_pred=predicted,
      ↪labels=['good', 'bad'])
```

```
[9]: array([[23092, 1386],
            [ 11, 99017]], dtype=int64)
```

```
[10]: metrics.plot_confusion_matrix(dt, test.loc[:, pred_vars], test['label'],
      ↪labels=['good', 'bad'])
      plt.show()
```



```
[11]: # compute baseline accuracy (predict all bad)
baseline = test_labels_stats['bad'] / (test_labels_stats['good'] +
    ↳ test_labels_stats['bad'])
print("Baseline accuracy is:", baseline)

# compute the observed accuracy
acc = metrics.accuracy_score(test['label'], predicted)
print("Observed accuracy is:", acc)
```

Baseline accuracy is: 0.8018071996502194
Observed accuracy is: 0.9886888086408757

```
[12]: result = metrics.classification_report(test['label'], predicted, digits=4)
print(result)
```

	precision	recall	f1-score	support
bad	0.9862	0.9999	0.9930	99028
good	0.9995	0.9434	0.9706	24478
accuracy			0.9887	123506
macro avg	0.9929	0.9716	0.9818	123506
weighted avg	0.9888	0.9887	0.9886	123506

1 Exercises

Try with different predictor variables. Does the model improve?

Try with different parameters for the tree. The list of adjustable parameters is here (Optional): Try running the models with the full dataset.

2 Task 1

3 models a base, a improvement one, and changing parameters.

```
[13]: # define new tree
dt_2 = tree.DecisionTreeClassifier()
# train the model using a list of column names
pred_vars_2 = ['count', 'num_root']
# The value we are trying to predict is 'label'
dt_2.fit(train.loc[:, pred_vars_2], train['label'])
# tree.plot_tree(dt)
```

```
[13]: DecisionTreeClassifier()
```

```
[14]: train, test = train_test_split(netattacks, test_size=0.25)
print("Rows in train:", len(train))
print("Rows in test:", len(test))
```

Rows in train: 370515

Rows in test: 123506

```
[15]: predicted_2 = dt_2.predict(test.loc[:, pred_vars_2])
print(predicted_2[:5]) # show first five predictions
```

```
['bad' 'bad' 'bad' 'good' 'bad']
```

```
[16]: from collections import Counter
# count test data
test_labels_stats_2 = Counter(test['label'])
print("Labels in the test data:", test_labels_stats_2)

# count predicted
predicted_labels_stats_2 = Counter(predicted_2)
print("Labels in the predictions:", predicted_labels_stats_2)
```


Labels in the test data: Counter({'bad': 99223, 'good': 24283})

Labels in the predictions: Counter({'bad': 97446, 'good': 26060})

```
[17]: metrics.confusion_matrix(y_true=test['label'], y_pred=predicted_2,
    ↪ labels=['good', 'bad'])
```

```
[17]: array([[23954,   329],
       [ 2106, 97117]], dtype=int64)
```

```
[18]: metrics.plot_confusion_matrix(dt_2, test.loc[:, pred_vars_2], test['label'],
    ↪ labels=['good', 'bad'])
plt.show()
```



```
[19]: # compute baseline accuracy (predict all bad)
baseline = test_labels_stats_2['bad'] / (test_labels_stats_2['good'] +
    ↪ test_labels_stats_2['bad'])
print("Baseline accuracy is:", baseline)

# compute the observed accuracy
acc = metrics.accuracy_score(test['label'], predicted_2)
print("Observed accuracy is:", acc)
```

Baseline accuracy is: 0.8033860703123735

Observed accuracy is: 0.9802843586546403

```
[20]: result = metrics.classification_report(test['label'], predicted_2, digits=4)
print(result)
```

	precision	recall	f1-score	support
bad	0.9966	0.9788	0.9876	99223
good	0.9192	0.9865	0.9516	24283
accuracy			0.9803	123506

macro avg	0.9579	0.9826	0.9696	123506
weighted avg	0.9814	0.9803	0.9805	123506

2.0.1 Improvement

```
[22]: #for loop to add one to see if I can beat 98.9
var_list = ['wrong_fragment', 'urgent', 'hot', 'num_failed_logins',
↳ 'num_compromised',
           'num_file_creations', 'num_shells', 'num_access_files',
↳ 'num_outbound_cmds',
           'error_rate', 'rerror_rate', 'same_srv_rate', 'diff_srv_rate',
↳ 'srv_count',
           'srv_error_rate', 'srv_rerror_rate', 'srv_diff_host_rate']

from collections import Counter

for i in var_list:
    dt_loop = tree.DecisionTreeClassifier()
    pred_vars_loop = ['count', 'num_root', i]
    dt_loop.fit(train.loc[:, pred_vars_loop], train['label'])
    train, test = train_test_split(netattacks, test_size=0.25)
    predicted_loop = dt_loop.predict(test.loc[:, pred_vars_loop])
    test_labels_stats_loop = Counter(test['label'])
    predicted_labels_stats_loop = Counter(predicted_loop)
    acc = metrics.accuracy_score(test['label'], predicted_loop)
    if acc > .9803:
        print(pred_vars_loop, ' Greater than first run: ', acc)

    if acc > .989:
        print(pred_vars_loop, ' Greater than professor run!')
```

```
['count', 'num_root', 'wrong_fragment'] Greater than first run:
0.9811588100982949
['count', 'num_root', 'hot'] Greater than first run: 0.9850938415947403
['count', 'num_root', 'num_compromised'] Greater than first run:
0.9838226482923906
['count', 'num_root', 'error_rate'] Greater than first run:
0.9838793257007756
['count', 'num_root', 'rerror_rate'] Greater than first run:
0.9808835198289961
['count', 'num_root', 'same_srv_rate'] Greater than first run:
0.9841222288795686
['count', 'num_root', 'diff_srv_rate'] Greater than first run:
0.9840088740627986
['count', 'num_root', 'srv_count'] Greater than first run: 0.9853367447735333
['count', 'num_root', 'srv_error_rate'] Greater than first run:
```

```
0.9846566158729131
['count', 'num_root', 'srv_error_rate'] Greater than first run:
0.9808187456479848
['count', 'num_root', 'srv_diff_host_rate'] Greater than first run:
0.9809644875552604
```

```
[23]: var_list = ['wrong_fragment', 'urgent', 'hot', 'num_failed_logins',
    → 'num_compromised',
    → 'num_file_creations', 'num_shells', 'num_access_files',
    → 'num_outbound_cmds',
    → 'serror_rate', 'rerror_rate', 'same_srv_rate', 'diff_srv_rate',
    → 'srv_error_rate', 'srv_rerror_rate', 'srv_diff_host_rate']

for i in var_list:
    dt_loop = tree.DecisionTreeClassifier()
    pred_vars_loop = ['count', 'num_root', 'srv_count', i]
    dt_loop.fit(train.loc[:, pred_vars_loop], train['label'])
    train, test = train_test_split(netattacks, test_size=0.25)
    predicted_loop = dt_loop.predict(test.loc[:, pred_vars_loop])
    test_labels_stats_loop = Counter(test['label'])
    predicted_labels_stats_loop = Counter(predicted_loop)
    acc = metrics.accuracy_score(test['label'], predicted_loop)
    if acc > .9803:
        print(pred_vars_loop, ' Greater than first run: ', acc)

    if acc > .989:
        print(pred_vars_loop, ' Greater than professor run!')
```

```
['count', 'num_root', 'srv_count', 'wrong_fragment'] Greater than first run:
0.9871504218418539
['count', 'num_root', 'srv_count', 'urgent'] Greater than first run:
0.9856363253607112
['count', 'num_root', 'srv_count', 'hot'] Greater than first run:
0.9905591631175813
['count', 'num_root', 'srv_count', 'hot'] Greater than professor run!
['count', 'num_root', 'srv_count', 'num_failed_logins'] Greater than first run:
0.985425809272424
['count', 'num_root', 'srv_count', 'num_compromised'] Greater than first run:
0.9899761954884783
['count', 'num_root', 'srv_count', 'num_compromised'] Greater than professor
run!
['count', 'num_root', 'srv_count', 'num_file_creations'] Greater than first
run: 0.985782067267987
['count', 'num_root', 'srv_count', 'num_shells'] Greater than first run:
0.9855310673165676
['count', 'num_root', 'srv_count', 'num_access_files'] Greater than first run:
0.985199099638884
['count', 'num_root', 'srv_count', 'num_outbound_cmds'] Greater than first run:
```

```

0.9859035188573835
['count', 'num_root', 'srv_count', 'serror_rate'] Greater than first run:
0.9862597768529464
['count', 'num_root', 'srv_count', 'rerror_rate'] Greater than first run:
0.9861383252635499
['count', 'num_root', 'srv_count', 'same_srv_rate'] Greater than first run:
0.9856768092238434
['count', 'num_root', 'srv_count', 'diff_srv_rate'] Greater than first run:
0.9854581963629298
['count', 'num_root', 'srv_count', 'srv_serror_rate'] Greater than first run:
0.9868832283451816
['count', 'num_root', 'srv_count', 'srv_rerror_rate'] Greater than first run:
0.9864864864864865
['count', 'num_root', 'srv_count', 'srv_diff_host_rate'] Greater than first
run: 0.9868265509367966

```

```


[27]: dt_loop = tree.DecisionTreeClassifier()
pred_vars_loop = ['count', 'num_root', 'srv_count', 'hot']
dt_loop.fit(train.loc[:, pred_vars_loop], train['label'])
train, test = train_test_split(netattacks, test_size=0.25)
predicted_loop = dt_loop.predict(test.loc[:, pred_vars_loop])
test_labels_stats_loop = Counter(test['label'])
metrics.plot_confusion_matrix(dt_loop, test.loc[:, pred_vars_loop],
    ↳test['label'], labels=['good', 'bad'])
predicted_labels_stats_loop = Counter(predicted_loop)
acc = metrics.accuracy_score(test['label'], predicted_loop)
print(acc)
plt.show()

```

```

0.9908587437047592


```

3 Task 2

```
[35]: dt_loop = tree.DecisionTreeClassifier(criterion='entropy', max_features='sqrt')
pred_vars_loop = ['count', 'num_root', 'srv_count', 'hot']
dt_loop.fit(train.loc[:, pred_vars_loop], train['label'])
train, test = train_test_split(netattacks, test_size=0.25)
predicted_loop = dt_loop.predict(test.loc[:, pred_vars_loop])
test_labels_stats_loop = Counter(test['label'])
metrics.plot_confusion_matrix(dt_loop, test.loc[:, pred_vars_loop],
    ↳test['label'], labels=['good', 'bad'])
predicted_labels_stats_loop = Counter(predicted_loop)
acc = metrics.accuracy_score(test['label'], predicted_loop)
print(acc)
plt.show()
```

0.9906482276164721




4 Optional Model on full Dataset

```
[39]: # use the column names
netattacksfull = pd.read_csv('data/kddcup.data.gz', names=name_list,
    ↳header=None, index_col=None)
netattacksfull['label'] = np.where(netattacksfull['type'] == 'normal.', 'good',
    ↳'bad')
```

```
[41]: dt_loop = tree.DecisionTreeClassifier(criterion='entropy', max_features='sqrt')
pred_vars_loop = ['count', 'num_root', 'srv_count', 'hot']
train, test = train_test_split(netattacksfull, test_size=0.25)
dt_loop.fit(train.loc[:, pred_vars_loop], train['label'])
predicted_loop = dt_loop.predict(test.loc[:, pred_vars_loop])
test_labels_stats_loop = Counter(test['label'])
metrics.plot_confusion_matrix(dt_loop, test.loc[:, pred_vars_loop],
    ↳test['label'], labels=['good', 'bad'])
predicted_labels_stats_loop = Counter(predicted_loop)
acc = metrics.accuracy_score(test['label'], predicted_loop)
print(acc)
plt.show()
```

0.993478729519977



[]: